

# Milestone 1: Team Formation & Architecture (Design Only)

**Project Title:** Deepfake & Synthetic Media Detector

**Team Name:** WCU RAMS

**Team Members:** Kevin Ha, Isreal Adegbie, Thomas Burke,  
Mo Abdularazzak, Alym Rejepov

**Date:** September 9, 2025

## 1 Overview

The **Deepfake & Synthetic Media Detector** is a Kubernetes-based pipeline designed to help users detect AI-generated or manipulated media, including images, videos, audio, and text. Users upload media, which is processed through a multi-pod pipeline to classify content as “Real” or “Fake,” providing confidence scores and logging results. This project emphasizes infrastructure as code, multi-pod orchestration, and persistence handling in Kubernetes.

### Problem

Misinformation from AI-generated or manipulated media is growing rapidly; users have no simple way to verify authenticity across text, audio, image, and video.

### Solution

Users upload text, videos, images, or audio. The pipeline analyzes authenticity using anomaly detection and compares it against trusted sources, flagging suspected deepfakes.

## 2 Use Cases & Workflow

### User Stories:

- As a user, I would upload my file (text, audio, image, or video) and receive a report within seconds stating if it is a deepfake or not.
- As a journalist, I would want to double-check my story’s origins and its authenticity.
- As an older person or someone who is not technically advanced, I would want my phone calls, emails, texts, etc., to be checked to see if they are robocalls or scams.
- As a teacher, I would want to upload my students’ assignments and check if they are AI-generated or not.
- As law enforcement, I would want to validate video/audio evidence or alibi authenticity before trial.

### 3 System Diagrams

#### Data-Flow Diagram

#### Deepfake & Synthetic Media Detector

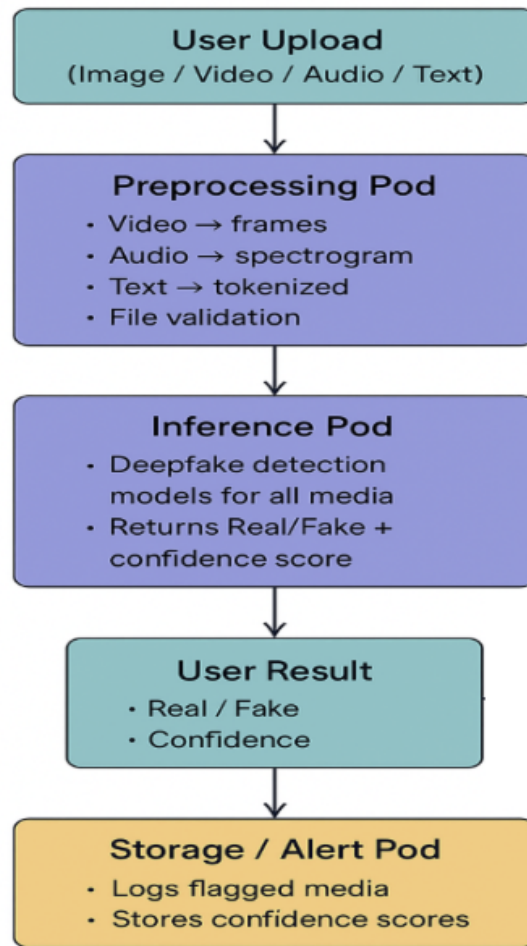


Figure 1: Data Flow Diagram of Media Processing Pipeline

## Kubernetes System Architecture




Figure 2: Kubernetes System Architecture for Deepfake Detection

## 4 Kubernetes Architecture Design

### Pod Structure:

- **API Gateway (Single Pod):** Handles all incoming requests from users. Kept separate to isolate the internet-facing component. Exposed externally via a LoadBalancer/Ingress with HTTPS.
- **Preprocessing (Single Pod):** Cleans and normalizes uploaded media (images, audio, video, text). Sandboxes raw, untrusted user data.
- **Inference (Multi-container Pod):** Runs the detection model with a metrics sidecar. Scales independently under load.
- **Storage & Alerts (Single Pod + Persistent Volume):** Stores results, logs, user credentials, and flagged media. Encryption at rest enabled.
- **User Management (Single Pod):** Handles sign-up/login; stores hashed credentials in Kubernetes Secrets.

### Service Types:

- **ClusterIP:** Internal communication between pods.
- **LoadBalancer:** External secure access for users.

## 5 API Contracts

### POST /auth/register

```
Input: { "username": "alice", "password": "mypassword123" }
Output: { "status": "success", "message": "Account created" }
```

### POST /auth/login

```
Input: { "username": "alice", "password": "mypassword123" }
Output: { "token": "<JWT-token>", "expires": "3600s" }
```

## POST /upload

```
Auth: Bearer JWT required
Input: multipart form-data { file: <media>, user_id: <uuid> }
Output: { "job_id": "12345", "status": "processing" }
```

## GET /result/job<sub>*i*</sub>*d*

```
Auth: Bearer JWT required
Output: { "status": "complete", "verdict": "fake", "confidence": 0.94 }
```

## 6 Authentication & Authorization

Authentication is required before any media upload or result retrieval. Users must first create an account with a username and password, which are hashed and stored in Kubernetes Secrets.

- **Normal Users** – can upload media and view their own results.
- **Admins** – can view all flagged content, system metrics, and manage users.

## 7 Risks

- Model size may cause slow startup.
- Large input files may exceed memory limits.
- Some file types may be unsupported.

## 8 Test Matrix

| Scenario         | Input           | Expected Output                    |
|------------------|-----------------|------------------------------------|
| Happy Path       | Valid image     | Prediction + confidence score      |
| Large File       | Video >100MB    | Error: “File too large”            |
| Invalid Format   | .exe file       | Error: “Unsupported File Type”     |
| Adversarial Case | Corrupted audio | Report: “Unable to process” + logs |

## 9 Team Charter

### Roles

- Lead / PM – Kevin
- DevOps – Isreal
- Backend / Interface – Thomas (frontend/interface), Mo and Aym (backend)
- Preprocessing – Thomas
- QA Docs – Kevin

### Meeting Cadence

- Weekly on Wednesday at 1:00 PM (in person)
- Quick check-ins as needed (virtual)

### **Definition of Done**

- Inference pod scales under load.
- Logs stored successfully.
- API passes happy-path and error-path tests.
- All services deployed in Kubernetes cluster.
- Program works as planned.
- Program passes all test cases.